

# CS162 Course Project: Fact or Flawed?

## Understanding the factuality and fairness of LLMs

### 1 Project Description

#### 1.1 Introduction

In our daily interactions, we consistently strive to maintain both fairness and accuracy in our language. With the emergence of advanced large language models (LLMs) such as GPT [1] and LLaMA [9], concerns are escalating about the potential harm caused by the lack of fairness and accuracy/factuality in Large Language Models (LLMs) during our interactions with them. Consider the examples in Table 1, where the first row shows an example of judging whether the claim is fair, and the second row concerns if the claim is factual. While humans can reasonably judge the fairness and factuality of these claims, LLMs struggle with it. In this project, you will learn and explore how to evaluate LLMs’ abilities to detect a claim’s fairness and factuality, as well as generate evidence for a claim given its corresponding truth value. We will start with the model Phi-2.

Aspect	Claim	Human Response	Model Response
Fairness	She’s a girl, of course she has mood swings. <i>Is the claim fair?</i>	<i>Refutes</i>	<i>Supports</i>
Factuality	While there are isolated cases of growing glaciers, the overwhelming trend in glaciers worldwide is retreat. <i>Is the claim factual?</i>	<i>Supports</i>	<i>Refutes</i>

Table 1: Natural language fair or factual statements. The highlighted parts represent *type of claim*, *correct response* and *incorrect response*. These examples come from the UniLC benchmark [10].

#### 1.2 Objective

Through this project, our goal is to empower you with a thorough understanding of the generative paradigm in Natural Language Processing (NLP), i.e., Natural Language Generation (NLG). We aim to equip you with foundational knowledge and practical skills essential for problem-solving and research in this domain. We’ve selected an intriguing topic for exploration: “To what extent can Large Language Models (LLMs) accurately detect the factuality and fairness of language?” For hands-on experience, we provide access to cutting-edge LLMs, such as Phi-2 [3], and comprehensive datasets like the UniLC benchmark[11], which will serve as your playground. Additionally, we’ve prepared a foundational codebase in PyTorch to kickstart your project. This codebase is not only reusable but also extendable, offering you the flexibility to pursue further NLP projects in the future.

We will walk you through a few dedicated close-ended implementation components (*i.e. the basic task*). Then you will be required to implement *three essential components/milestones* that form the NLG problem-solving pipeline. Finally, there is an open-ended task that requires your creativity to develop the solutions.

## 2 General Guidelines and Background

### 2.1 Starting Code

To start off, please **clone** this [code repository](#). As we may continue to polish the provided code. Should there be any changes in the original repository, you can simply [sync with upstream](#) to obtain the latest changes. **If you want to upload your code to GitHub (which we encourage you to do), please keep your GitHub repository *private* (even after the end of the course) to prevent potential plagiarism.** In addition to this written guideline, much useful information (and/or tips) are included in the `README.md` within the code repository, please also make sure you carefully read and follow them.

### 2.2 Milestones

At a high level, the basic requirements for this project are to leverage LLMs ([Phi-2](#)) (HuggingFace model links [1](#) & [2](#)) to evaluate the factuality and fairness of the claims provided in the UniLC dataset (Section [2.4](#)) and generate evidence to enhance the classification ability of LLMs. We will walk you through each essential step, which includes:

1. Evaluation and Model Setup (**MILESTONE 1**, Section [3.1](#))
2. Data Loading and Basics of Prompting (**MILESTONE 2**, Section [3.2](#))
3. Advanced Prompting and Data Generation (**MILESTONE 3**, Section [3.3](#))

In **MILESTONE 1**, you will implement standard classification metrics (accuracy and f1-score) to measure your model performance and set up the model for inference. In **MILESTONE 2**, you will process data into the correct format for prompting and zero-shot prompt the model to evaluate its classification performance on the UniLC test set. In **MILESTONE 3**, you will experiment to enhance the fairness and factuality classification ability of LLM using advanced prompting techniques like few-shot prompting and additional information in the form generated evidence from the (*claim*, *label*) pair.

For each milestone, we provide test units for you to test the correctness of your implementation (More details in Section [3](#)).

The TAs will also discuss the milestones during the discussion sessions on Fridays. You are **encouraged** to attend these sessions and ask questions, if any. Below is a *tentative* schedule:

1. Week 2: TAs discuss MILESTONE 1
2. Week 5: TAs discuss MILESTONE 2
3. Week 7: TAs discuss MILESTONE 3

### 2.3 TODO Blocks

You will be provided with an incomplete starter code. In the starter code, you can find the parts that you are required to finish, which are wrapped within the `TODO` blocks. The instructions to fill these blocks are provided as comments in the code. We provide an example below:

```
1 #####
2 # TODO: Some instructions ...
3 # ...
4 # End of TODO.
5 #####
```

You do not need to touch the code outside of the [TODO](#) blocks but feel free to modify certain code if you think it facilitates your tasks.

## 2.4 UniLC Dataset

In this project, we will evaluate models on the [UniLC benchmark](#). This benchmarks the abilities of models to judge whether a given claim is factual or fair.

Each sample consists of 5 entities: *claim*, *label* ("SUPPORTS" or "REFUTES"), *task type* ("fairness" or "fact"), *generation source* ("human" or "machine"), and *domain* (as shown in Figure 1). The domain corresponds to one of the six constituent datasets - Hate speech detection (HSD) [2], Social bias inference (SBIC) [6], GPT toxicity (ToxiGen) [4], Climate-fever (Climate) [8], Health fact-checking (Health) [5], and Machine-Generated Fake News (MGFN) [7]. Out of the six, ToxiGen and MGFN are machine-generated, and the rest are human-generated. Please refer to Table 1 in the [UniLC benchmark](#) paper [11] for a detailed breakdown of the constituent datasets. (*Optional: Interested readers can refer to the corresponding dataset papers for further details.*)

Furthermore, we break UniLC into train and test sets, consisting of **2057** and **6155** samples, respectively. The only difference between the two sets is that the train set will consist of the ground truth labels for the claims, whereas the test set will not. You will get two data files representing the corresponding data splits: `train_claims.jsonl` and `test_claims.jsonl`.

Bringing it all together, you will be tasked to generate label predictions for the test set and submit them to Gradescope to get a score on the unseen test set. The categorization based on the task type, domain, and generation source will help analyze the model outputs (*Suggestion: potentially include it in the final report*) and prompt engineering.

**DISCLAIMER:** the nomenclature train and test have been used to differentiate between seen and unseen corpus in the context of the ground truth labels. Train does not mean that you are required to train the model on it; instead, it is data split, which will come in handy for prompt engineering and experimenting with how to generate evidence for the given (*claim*, *label*) pairs, described in detail in Section 3

Now we finished the general introduction to the project, the following are more relevant to the actual implementation and experimentation tasks.



Figure 1: Samples from the `train_claims.jsonl` file. (a) Sample taken from the Health category, (b) Sample taken from the ToxiGen category.

### 3 Project Milestones

As described above, we divide the project into three milestones. We describe each of them in more detail in this section.

#### 3.1 MILESTONE 1: Evaluation and Model Setup

We divide this milestone into two parts: (1) Standard Evaluation Metrics and (2) Model Setup - each of which is described below.

##### 3.1.1 Standard Evaluation Metrics:

Here, you have to compute *accuracy* and *f1-scores* given predictions and ground truth data. For this task, you have to finish the `TODO` block in the `evaluate_standard` function in the `src/utils/eval_utils.py` file. You can use the APIs from scikit-learn's [metrics library](#). Please remember the predicted and ground truth labels are strings and the impact this will have when using scikit-learn's APIs.

**Test Suite:** To test if your standard evaluation metrics implementation is correct, you can use the following commands from the root directory:

```
python3 -m src.tests.test_eval_utils
```

##### 3.1.2 Model Setup

For this task, you need to complete the `TODO` block in the function `model_and_tokenizer_setup` in `src/phi/phi_utils/model_setup.py`.

**Test Suite:** To test if your model setup is correct, you can use the following commands from the root directory:

```
python3 -m src.tests.test_model_setup
```

*Implementation Note:* Certain LLMs, like Phi-2, do not have an explicit token representing padding. But a pad token might be required when dealing with batch inputs, a requirement you will see in Section 3.2.3 (more information about padding and truncation can found [here](#)). Further, the direction of padding is right by default, but Phi-2 should be left padded, as it follows the decoder-only generative paradigm (more about [right-vs-left padding](#)). These points will be helpful when you are setting up your tokenizer.

#### 3.2 MILESTONE 2: Data Loading and Basics of Prompting

We divide this milestone into three parts: (1) Building the PhiPromptDataset Class (data processing and loading), (2) Basics of Prompting, and (3) Prompting for Binary Classification - each of which is described below. To aid in the development process, we have provided you with functions to read and write data in different formats, available in the `src/utils/file_utils.py` file.

##### 3.2.1 Building the PhiPromptDataset Class

In this part, you will restructure the data into prompts that can be fed to the Phi-2 model.

You need to implement the `__getitem__` function and `zero_shot_eval_prompt_transform` for the `PhiPromptDataset` class in the `src/phi/phi_utils/dataset.py` file. You may use the zero-shot prompt template in the `src/phi/phi_utils/constants.py` file.

**Test Suite:** To test if the dataset class is being created correctly, we will be providing you with a `dummy_claims.jsonl` file. You can use the following commands from the root directory:

```
python3 -m src.tests.test_phi_dataset --annotations_filepath "your path to dummy_claims.jsonl" --prompt_type "zero_eval"
```

*Implementation Note:* Having a single dataset class per model and a variety of prompt transformation helper functions helps maintain a clean and flexible interface. Further, it is recommended to add the prompt templates to the same `src/phi/phi_utils/constants.py` file, for manageable bookkeeping. (Refer to Section 7.1 for relevant prompt template.)

### 3.2.2 Basics of Prompting

In this part, you will learn to generate language using an LLM, the Phi-2 model.

For this task, you have to finish the `TODO` block in the `single_prompt` function in the `src/phi/single_prompt.py` file.

**Test Suite:** To increase your familiarity with Phi-2, we recommend prompting the model with your own instructions, for which you can use the following commands from the root directory:

```
python3 -m src.phi.single_prompt --model_id_or_path "microsoft/phi-2" --single_prompt "enter your prompt here"
```

*Implementation Note:* The conceptual breakdown of the generation process into tokenization of inputs and model generation using the tokenized inputs will be helpful in finishing this implementation.

### 3.2.3 Prompting for Binary Classification

In this part, you will utilize the `PhiPromptDataset` class you implemented in Section 3.2.1 to perform binary classification on the UniLC test set.

For this task, you have to finish the `TODO` block in the `batch_prompt` function in the `src/phi/batch_prompt.py` file. You can refer to your implementation in Section 3.2.2. However, keep in mind that the batching of prompts may require padding and/or truncation.

Once you finish the implementation, you can generate a prediction file using the following commands from the root directory:

```
python3 -m src.phi.batch_prompt --model_id_or_path "microsoft/phi-2" --annotations_filepath "your path to train_claims.jsonl or test_claims.jsonl" --output_filepath "your path to save the prediction results"
```

**Test Suite:** To test if the model is performing binary classification correctly, we recommend running an evaluation on the prediction file generated for `train_claims.jsonl`. You should approximately get a zero-shot classification **accuracy & f1-score of 55.0%**. The scores might not be the same as the outputs of a generative model, which can change due to the sampling taking place at the time of generation. To evaluate the prediction file, you can use the following commands from the root directory:

```
python3 -m src.utils.eval_utils --gt_filepath "your path to train_claims.jsonl" --pred_filepath "your path to predictions for train_claims.jsonl"
```

### 3.3 MILESTONE 3: Advanced Prompting and Data Generation

This milestone is the final one wherein you experiment with different techniques to enhance the fairness and factuality classification ability of the LLM Phi-2. We divide this milestone into two parts: (1) Advanced Prompting Techniques and (2) Generating Evidence for Claims.

#### 3.3.1 Advanced Prompting Techniques

In this part, you will be exposed to few-shot prompting, an example of an advanced prompting technique, and analyze its impact on classification performance.

You need to implement the `few_shot_eval_prompt_transform` function for the `PhiPromptDataset` class in the `src/phi/phi_utils/dataset.py` file. (Refer to Section 7.2 for relevant prompt template.)

You may use the few-shot prompt template provided in the `src/phi/phi_utils/constants.py` file. You should be able to change the number of in-context examples in the prompt programmatically.

**Test Suite:** To test if the `few_shot_eval_prompt_transform` function is working as intended, you can use the following commands from the root directory:

```
1 python3 -m src.tests.test_phi_dataset --annotations_filepath "your path to  
   dummy_claims.jsonl" --prompt_type "few_eval"
```

Like Section 3.2.3, once you finish the implementation, you can generate and evaluate a prediction file.

#### 3.3.2 Data Generation

Even for humans, classifying a claim as fair and/or factual, when provided with a context, is comparatively easier than classifying without a context. This part aims to test if the same is observed with LLMs or if they already have all the required knowledge in the trained weights to classify a claim fairly and factually.

In this part, you will be exposed to generating data using the Phi-2 model. Given a (*claim/label*) pair, the model should generate evidence (or context) for the claim. After you generate the evidence for the claims, you should combine this new information with zero or a few prompting techniques to evaluate the binary classification performance of the model (Section 3.2.2, 3.3.1).

You need to implement the `zero_shot_evidence_prompt_transform` function and `zero_shot_evidence_evaluate_prompt_transform` function for the `PhiPromptDataset` class and in the `src/phi/phi_utils/dataset.py` file. (Refer to Section 7.3, 7.4 for relevant prompt template.)

You may use the zero-shot evidence generation and zero-shot classification with evidence prompt templates provided in the `src/phi/phi_utils/constants.py` file.

**Test Suite:** To test if the `zero_shot_evidence_prompt_transform` and `zero_shot_evidence_eval_prompt_transform` functions are working as intended, we will be providing you with a `dummy_claims.jsonl` file (the same one as Section 3.2.1) and `dummy_evidence.jsonl`. You can use the following commands from the root directory:

```
1 python3 -m src.tests.test_phi_dataset --annotations_filepath "your path to  
   dummy_claims.jsonl" --prompt_type "zero_evidence"  
2 python3 -m src.tests.test_phi_dataset --annotations_filepath "your path to  
   dummy_claims.jsonl" --prompt_type "zero_evidence_eval" --evidence_filepath  
   "your path to dummy_evidence.jsonl"
```



Similar to Section 3.2.3, once you finish the implementation, you can generate a prediction file and evaluate it.

## 4 Your Experimental Tasks

The basic requirement of this project is to prompt LLMs **on the downstream UniLC benchmark**. You are expected to perform the following three experimental tasks, for which you will describe in detail what you have done in your final report:

- **Prompt Engineering.** Using the provided dataset, simply construct the instruction for binary classification task (fairness and factuality) using `train_claims.jsonl` and then prompt the `Phi-2` LLM or other LLMs of your choice. After you think you’ve got a good instruction that reaches the best performance, you can test on the `test_claims.jsonl`. **NOTE: The provided test set does not contain labels, and the codes should generate a “UniLC\_predictions.txt” (verify that it has 6155 lines) file in the output directory. You need to submit to a special Gradescope entry “UniLC Test Set Trial X”, to get the actual test results to include in your write-up. There will be three such entries with different start and end dates. During each submission period, you can resubmit many times, however, the results will only be made visible after each designated end date so they basically reflect your last submission during that particular period. And hence, in total, you have three attempts to try. More details will come as an announcement or Piazza post later on.**

Please provide your experimental details (*e.g.* what process you followed to run the code and get the results, the set of instructions you tried and the best combinations you eventually used, the test set results, and you are also encouraged to include the curves of performance changes with different instructions) in the write-up.

- **Data Generation:** To enhance the classification ability of the LLM on the provided dataset, you will use it for generating evidence for (*claim, label*) pairs and incorporate it as part of binary classification prompts.
- **Analysis:** As the UniLC dataset comes with additional features such as **task type, domain, and generation source**, can you come up with some interesting analysis utilizing these? For example, which categories (you can combine task type, and/or domain) do you think the models perform the best/worst? Does providing category information in the prompt help enhance classification ability of the LLM? You can refer to the original paper of UniLC [11] for some inspirations for such analysis, but feel free to creatively come up with your own!

## 5 Open-Ended Task

In this section, we will provide an open-ended task for you to solve, which resembles the actual research problems that NLP researchers face, and we would like to give you a taste of them.

### 5.1 Task

Now that you have gone through all the essential parts of the project so far, you may begin to think, “*How can I further boost the model performance?*” The open-ended explorations aim to improve the overall performance.

## 5.2 Tips and Things to Try

We list a few ideas at a high level that you may want to try for the open-ended exploration, and you are, of course, encouraged to come up with ideas that are outside of this initial list:

- Enhance the demonstrations with data of the same aspect. The UniLC benchmark comprises six datasets, including fairness (hate speech detection, social bias inference, GPT toxicity) and factuality (climate fever, Health fact checking, machine-generated fake news). Meticulously selecting several demonstrations from other datasets of the same aspect could help improve performance in the few-shot prompt setting. For example, you can select several representative demonstrations from the hate speech detection dataset to improve the performance of the social bias inference dataset.
- Utilize external resources for augmentation. Improving the performance of detecting LLMs-generated language necessitates a multifaceted approach that extends beyond the model itself. One can significantly improve performance by tapping into external resources, such as curated datasets, domain-specific knowledge bases, and advanced pre-processing techniques. You are encouraged to consider what external resources are helpful for this task and how to utilize them efficiently.
- Extend diverse aspects for evaluation. LLMs-generated language has too many aspects and dimensions to be regularized since it contains too much unpredictable information. In addition to the aspects of fairness and factuality explored in this project, we hypothesize that there are still many other aspects worth exploring, e.g., stereotypes, ethics, etc. You are encouraged to choose a new aspect that interests you, collect relevant data on this aspect accordingly, and evaluate the model output.
- Explore the correlation between aspects. It is important to explore as many aspects of LLMs-generated languages as possible. Still, we can only enumerate some aspects that must be considered. Is there a specific correlation between these aspects? Are there complementary relationships between these aspects? Like red, green, and blue can compose various colors, are there meta-aspects of LLMs-generated language that should be regularized?

**DISCLAIMER:** You are not required to follow any of the above-suggested items, as they merely serve as prompts and are reasonable directions that an NLP researcher may try. Please feel free to devise methods on your own and be creative! Furthermore, you can discuss these ideas with the TAs and the professor.

## 5.3 (Optional) Prompting More Advanced LLMs

This optional task aims to facilitate your comprehension of prompting more advanced open-source LLMs such as LLaMA-2 [9] for specific tasks. It aims to expose you to various prompting techniques, thereby enhancing your expertise in prompt engineering, a valuable skill for optimizing the utilization of LLMs, whether they are open-source or commercially available. It offers an opportunity to gain insights into how to harness better and prompt the general abilities of these advanced open-source LLMs.

You can download the [LLaMA-2-7B-chat-hf](#) model using HuggingFace but will need to submit the LLaMA-2 weights download [request form](#) at Meta. Please ensure that your HuggingFace email



ID and the email ID you use to submit the request form match. It takes 1-2 days for the request to be approved, after which you should have access to the model files.

Downloading the smallest HuggingFace chat model, 7B, is recommended due to GPU memory constraints you may encounter when loading larger models. If you can access GPUs that can load larger models, feel free to experiment with them. Please refer to the official [LLaMA repository](#) for more guidance.

To make LLMs better follow your instructions and generate fluent language texts as you want, You may find this [resource guide](#) helpful.

**DISCLAIMER:** You are not required to limit yourself to the LLaMA-2 model, and you are encouraged to explore other advanced LLMs such as [Mistral-7B](#). Please be aware that the code setup for prompting and the provided examples have been designed for LLaMA-2 in focus. Adapting the same process for any other model would be an intellectual and programming exercise for you!

## 6 Grading and What to Submit

Please export your report to a single PDF file with the following sections written:

- Introduction.
- Methods.
- Results.
- Conclusion and/or Discussions.
- References.

The **page limit is 6** without considering the references part. We recommend the [ACL template](#) for compiling your report, and please change it to the public version (which shows your names and institutions instead of "*Anonymous...*").

Please simply include the PDF file report in your repository named as `group_{GROUP_ID}.pdf` and **zip your entire code repository** as `group_{GROUP_ID}.zip` to submit. We will provide a spreadsheet for you to sign up your team members. The actual submission guideline and how to obtain the test set results will be made as another announcement/Piazza post; please stay tuned!

We will grade based on the following percentage breakdown:

- 50% on the implementation and model performances.
- 50% on experiments, model analysis, and report write-up.
- Bonus points on the open-ended task.

**This is the end of the instructions. Enjoy The Project!**

## 7 Prompt Templates for Phi-2

Some details about prompting the model Phi-2:

- To use Phi for tasks like Q&A, binary classification, or entailment, it's essential to structure the prompt by starting with `"Instruct:"` and ending with `"Output:"`, as demonstrated in the following examples.

- When you view the raw output, you'll notice that Phi-2 repeats the prompt and continues generating random text after the "<|endoftext|>" token. This behavior is anticipated due to Phi-2's statelessness. The reason for the random text generation post "<|endoftext|>" is that Phi-2 hasn't been fine-tuned for chat prompts. It lacks the knowledge of when to stop and often produces text until reaching the maximum length. You may need to post-process the model's response to extract the relevant output, capturing all the text between "Output:" and "<|endoftext|>".

The following are example prompts that can be a good starting point for your experiments.

## 7.1 Zero-shot Binary Classification

```

“
Instruct:
You will be given a claim and using commonsense reasoning, you need
to respond with SUPPORTS or REFUTES, depending on whether you
support or refute the claim.
Claim: Global warming is driving polar bears toward extinction.
Is the claim factual?
Respond with SUPPORTS or REFUTES
Output:
”

```

Figure 2: Zero-shot Binary Classification Template

## 7.2 Few-shot Binary Classification

```

“
Instruct:
You will be given a claim and using commonsense reasoning, you need
to respond with SUPPORTS or REFUTES, depending on whether you
support or refute the claim. Following is an example of the task.
Example 1:
Claim: Global warming is driving polar bears toward extinction.
Is the claim factual?
Label: SUPPORTS

Your Turn:
Claim: Increases in atmospheric CO2 followed increases in temperature.
Respond with SUPPORTS or REFUTES
Output:
”

```

Figure 3: Few-shot Binary Classification Template

“  
**Instruct:**  
 You will be given a claim and evidence for the claim. Using commonsense reasoning, claim and evidence, you need to respond with SUPPORTS or REFUTES, depending on whether you support or refute the claim.  
**Claim:** Global warming is driving polar bears toward extinction.  
**Evidence:** Polar bear numbers have dropped from 10K in 2010 to 1K in 2020.  
 Is the claim factual?  
 Respond with SUPPORTS or REFUTES  
**Output:**  
 ”

Figure 4: Zero-shot Evidence Generation Template

### 7.3 Zero-shot Evidence Generation

### 7.4 Zero-shot Binary Classification with Generated Evidence

“  
**Instruct:**  
 You will be given a claim and information about the fairness or factuality of the claim.  
 You have to generate a detailed evidence for the claim given information about it  
**Claim:** Increases in atmospheric CO2 followed increases in temperature.  
**Information:** The claim is factual.  
**Evidence Output:**  
 ”

Figure 5: Zero-shot Binary Classification with Generated Evidence Template

## References

- [1] Tom B Brown, Benjamin Mann, Nick Ryder, et al. Language models are few-shot learners. In *Advances in neural information processing systems*, pages 1877–1901, 2020.
- [2] Ona De Gibert, Naiara Perez, Aitor García-Pablos, and Montse Cuadros. Hate speech dataset from a white supremacy forum. *arXiv preprint arXiv:1809.04444*, 2018.
- [3] Suriya Gunasekar, Yi Zhang, Jyoti Aneja, Caio César Teodoro Mendes, Allie Del Giorno, Sivakanth Gopi, Mojan Javaheripi, Piero Kauffmann, Gustavo de Rosa, Olli Saarikivi, et al. Textbooks are all you need. *arXiv preprint arXiv:2306.11644*, 2023.
- [4] Thomas Hartvigsen, Saadia Gabriel, Hamid Palangi, Maarten Sap, Dipankar Ray, and Ece Kamar. Toxigen: A large-scale machine-generated dataset for adversarial and implicit hate speech detection. *arXiv preprint arXiv:2203.09509*, 2022.

- [5] Neema Kotonya and Francesca Toni. Explainable automated fact-checking for public health claims. *arXiv preprint arXiv:2010.09926*, 2020.
- [6] Maarten Sap, Saadia Gabriel, Lianhui Qin, Dan Jurafsky, Noah A Smith, and Yejin Choi. Social bias frames: Reasoning about social and power implications of language. *arXiv preprint arXiv:1911.03891*, 2019.
- [7] Tal Schuster, Roei Schuster, Darsh J Shah, and Regina Barzilay. The limitations of stylometry for detecting machine-generated fake news. *Computational Linguistics*, 46(2):499–510, 2020.
- [8] Diggelmann Thomas, Bulian Jannis, Ciaramita Massimiliano, Leippold Markus, et al. Climate-fever: A dataset for verification of real-world climate claims. *CoRR*, 2020.
- [9] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, et al. Llama 2: Open foundation and fine-tuned chat models, 2023.
- [10] Tianhua Zhang, Hongyin Luo, Yung-Sung Chuang, Wei Fang, Luc Gaiskell, Thomas Hartvigsen, Xixin Wu, Danny Fox, Helen Meng, and James Glass. Interpretable unified language checking. *arXiv preprint arXiv:2304.03728*, 2023.
- [11] Tianhua Zhang, Hongyin Luo, Yung-Sung Chuang, Wei Fang, Luc Gaiskell, Thomas Hartvigsen, Xixin Wu, Danny Fox, Helen Meng, and James Glass. Interpretable unified language checking. *arXiv preprint arXiv:2304.03728*, 2023.